

IDS 572 Assignment 1 Part B

Lauren Sansone, Joshua Pollack, Lina Quiceno Bejarano

Variable Modifications: Remove loans with a status other than charged off and fully paid, changing emp_length to factor

5 Develop decision tree models to predict default.

(a) Split the data into training and validation sets. What proportions do you consider, why?

The Lending Club data was separated into three sets to build a decision tree model: training, validation and test to help build an accurate model. The training set was split into 70% of the data, using the majority of the data to train the model on predicting loan status. The validation set was split into 10% of the data, to cross evaluate the performance of the training data. The test set was split into 20% of the data, to test how well the model predicts loan status after training and validation. Set.seed set to 200 to generate the random number sequence.

5 (b) Train decision tree models (use both rpart, c50)

The decision tree model was trained using rpart and C50. Variables due to leakage were removed including debt settlement variables, hardship variables, payment amount variables, account balance variables, late fee variables, payment plan variables and others that would not be available at the time of the loan. The data set was cleaned, removing variables that could cause bias or were not necessary to predict loan status. Variables were removed if they had more than 60% of missing variables, including employment length and the X1 variable. Some variables with fewer missing values were replaced with median values such as months since recent inquiry. A derived attribute was created for the proportion of bank cards in satisfactory standing.

The rpart model was initially trained using the information index and min. split of 30. The model resulted in over fit, with very high accuracy and no charged off predictions. The same result was achieved when changing the model to the gini index. A complexity parameter was added, beginning with 0.0001. The decision tree had about an 85% accuracy but the tree was so large, it could not even be plotted – a result of extreme over fit. The cp values were experimented with and decreased, resulting in slightly smaller decision trees.

A summary of the training data displayed a total much larger number of fully paid loans than charged off loans. Since the totals were unbalanced, 3 to 1 weights were added to create a more balanced data set and encourage the model to predict charge offs.

A print out of the complexity parameters were given to find the best tree. The cp value within one standard deviation of the lowest x-error was determined to be 0.05 and incorporated in the training model to prune the tree. Before pruning, the training data predicted loan status with 80.7% accuracy and 88% specificity. After pruning, the accuracy increased to 81.4% and specificity of 90.7%. Since it was above 80%, the cross-validation data was run.

With the same parameters, the cross validation predicted 79.02% accuracy and 89.2% specificity. The test data was then run and predicted 77% accuracy and 90% specificity (precision). The performance of the tests was based on accuracy and specificity.

The performance of the model was evaluated using an ROC curve. The initial run of the curve had fairly poor separation and was mostly convex with small concavities. The curve was then lifted into 10 groups to evaluate performance.

Creating a Weighted Tree For the Training Set

```
myweights = ifelse(lcdfTrn$loan_status == "Charged Off", 3, 1 )

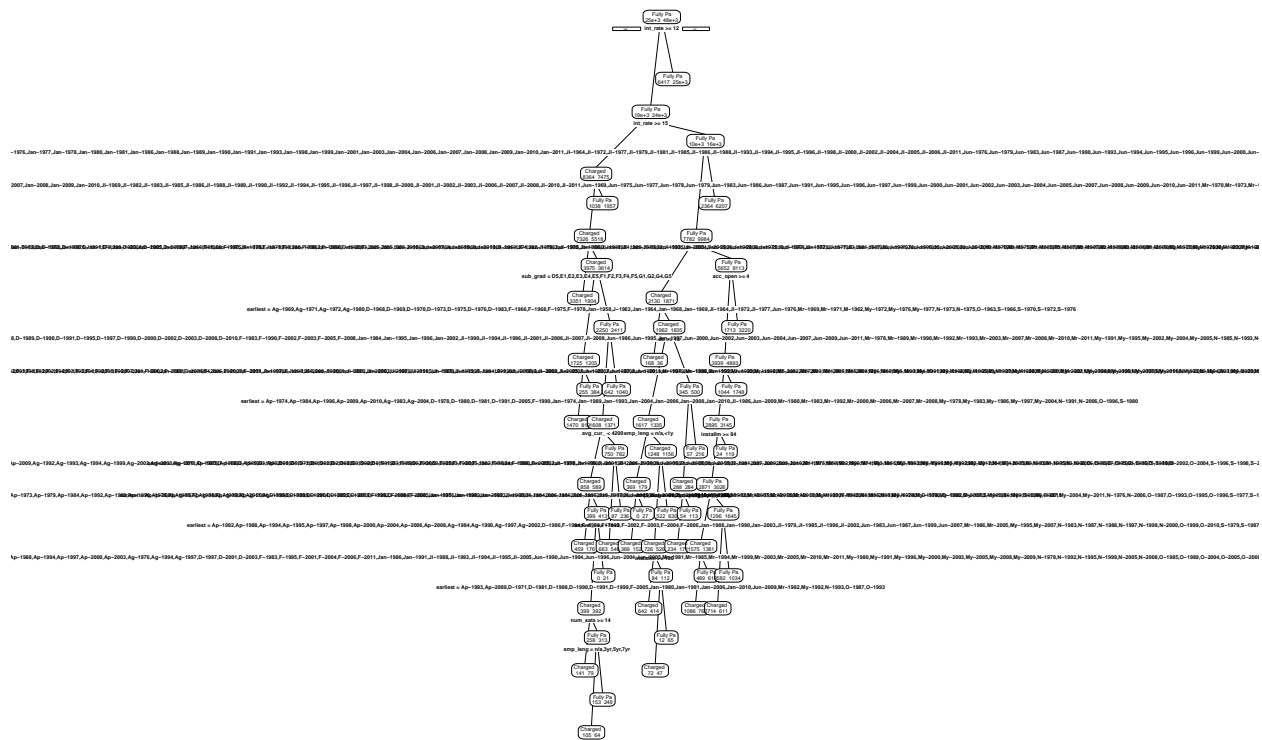
Wgtd_lcdT <- rpart(loan_status ~., data=lcdfTrn, method="class", weights = myweights, parms = list(spl

pred_wgtTrn=predict(Wgtd_lcdT,lcdfTrn, type='class')

#Confusion table
confusionMatrix(table(predWghtTrain = pred_wgtTrn, true=lcdfTrn$loan_status))

## Confusion Matrix and Statistics
##
##           true
## predWghtTrain Charged Off Fully Paid
##   Charged Off      3158      5781
##   Fully Paid       5151     42624
##
##           Accuracy : 0.8072
##           95% CI : (0.804, 0.8105)
##   No Information Rate : 0.8535
##   P-Value [Acc > NIR] : 1
##
##           Kappa : 0.2527
##
## Mcnemar's Test P-Value : 1.789e-09
##
##           Sensitivity : 0.38007
##           Specificity : 0.88057
##           Pos Pred Value : 0.35328
##           Neg Pred Value : 0.89218
##           Prevalence : 0.14651
##           Detection Rate : 0.05568
##   Detection Prevalence : 0.15762
##           Balanced Accuracy : 0.63032
##
##           'Positive' Class : Charged Off
##

Display the rpart Tree for training set
rpart.plot::prp(Wgtd_lcdT, type=2, extra=1)
```



Summary of lcDT

```
call:
rpart(formula = loan_status ~ ., data = lcdfTrn, weights = myweights,
      method = "class", parms = list(split = "information"), control =
rpart.control(cp = 0.001))
n= 56714
```

	CP	nsplit	rel error	xerror	xstd
1	0.024177264	0	1.0000000	1.0000000	0.005145924
2	0.005322208	3	0.9274682	0.9764914	0.005115765
3	0.005195170	6	0.9115016	1.0016849	0.005148022
4	0.005175111	8	0.9011112	1.0008826	0.005147024
5	0.003109079	9	0.8959361	1.0038512	0.005150708
6	0.002988727	11	0.8897180	1.0043327	0.005151303
7	0.002166326	13	0.8837405	1.0046937	0.005151749
8	0.001945681	15	0.8794079	1.0105508	0.005158927
9	0.001263690	21	0.8622779	1.0087857	0.005156774
10	0.001123280	23	0.8597505	1.0099089	0.005158145
11	0.001083163	24	0.8586272	1.0107514	0.005159171
12	0.001016301	25	0.8575440	1.0104305	0.005158780
13	0.001002929	29	0.8528503	1.0119549	0.005160633
14	0.001000000	30	0.8518474	1.0119549	0.005160633

Details About the Training Set


```
#tree size and performance for different complexity parameter values
printcp(Wghtd_lcDT)
```

```
##
## Classification tree:
```

```
## rpart(formula = loan_status ~ ., data = lcdfTrn, weights = myweights,
##       method = "class", parms = list(split = "information"), control = rpart.control(cp = 0.001))
##
## Variables actually used in tree construction:
## [1] acc_open_past_24mths annual_inc      avg_cur_bal
## [4] dti                earliest_cr_line  emp_length
## [7] installment        int_rate       num_sats
## [10] sub_grade
##
## Root node error: 24927/56714 = 0.43952
##
## n= 56714
##
##      CP nsplit rel error  xerror      xstd
## 1 0.0241773    0  1.00000 1.00000 0.0051459
## 2 0.0053222    3  0.92747 0.97649 0.0051158
## 3 0.0051952    6  0.91150 1.00168 0.0051480
## 4 0.0051751    8  0.90111 1.00088 0.0051470
## 5 0.0031091    9  0.89594 1.00385 0.0051507
## 6 0.0029887   11  0.88972 1.00433 0.0051513
## 7 0.0021663   13  0.88374 1.00469 0.0051517
## 8 0.0019457   15  0.87941 1.01055 0.0051589
## 9 0.0012637   21  0.86228 1.00879 0.0051568
## 10 0.0011233   23  0.85975 1.00991 0.0051581
## 11 0.0010832   24  0.85863 1.01075 0.0051592
## 12 0.0010163   25  0.85754 1.01043 0.0051588
## 13 0.0010029   29  0.85285 1.01195 0.0051606
## 14 0.0010000   30  0.85185 1.01195 0.0051606
```

#Plot

```
plotcp(Wghtd_lcDT)
```



#Variable importance as given by a decision tree model
Wghtd_lcDT\$variable.importance

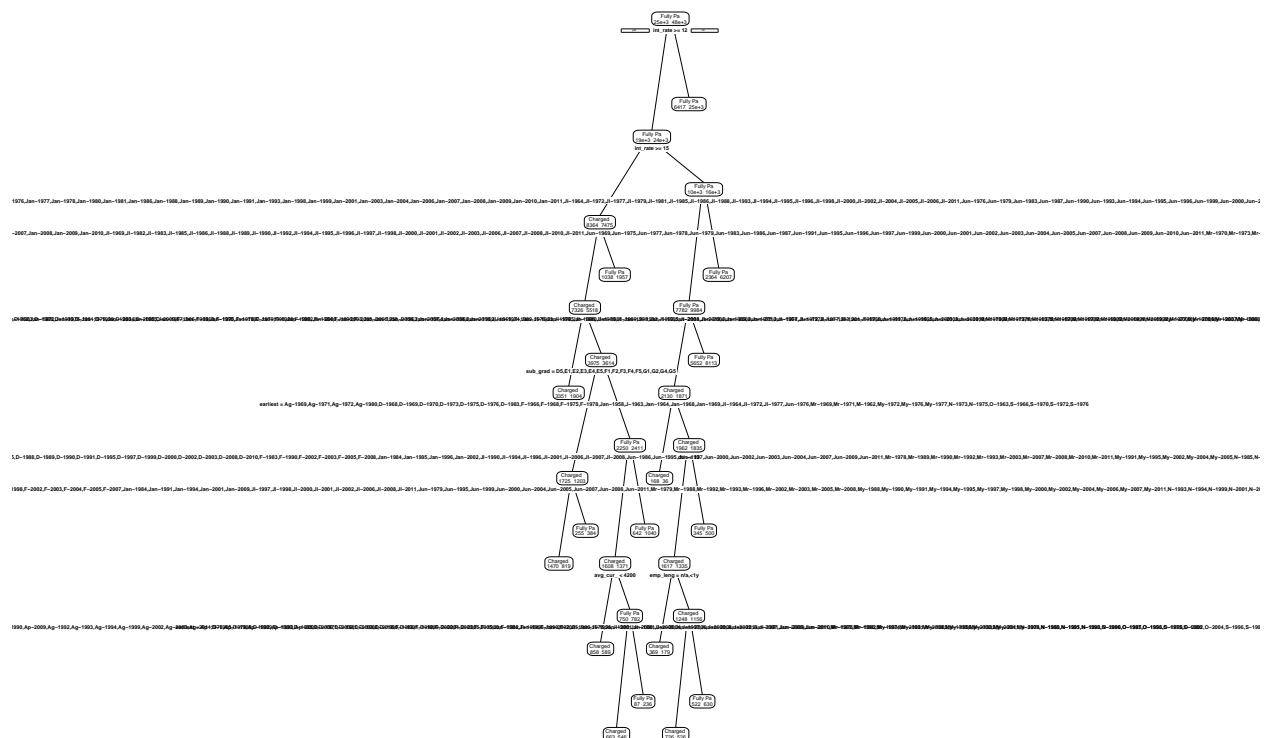
##	int_rate	sub_grade
##	2691.4752931	2287.6435343
##	grade	earliest_cr_line
##	1918.5165763	1224.5864403
##	bc_open_to_buy	total_bc_limit
##	473.4999214	405.5368413
##	total_rev_hi_lim	acc_open_past_24mths
##	331.4927267	64.8185512
##	installment	emp_length
##	64.3916895	57.8532145
##	mo_sin_old_rev_tl_op	avg_cur_bal
##	57.3582421	44.7476097
##	loan_amnt	tot_hi_cred_lim
##	43.4657633	37.9086366
##	dti	num_tl_op_past_12m
##	34.4490394	34.4139499
##	total_bal_ex_mort	num_sats
##	29.4843837	27.1732068
##	mort_acc	home_ownership
##	22.8458601	20.2000658
##	total_il_high_credit_limit	mo_sin_rcnt_rev_tl_op
##	19.9855626	17.1194350
##	annual_inc	mo_sin_rcnt_tl
##	16.6797016	14.9495017
##	mths_since_recent_bc	mo_sin_old_il_acct

```
##                13.7747058                12.1795358
##                purpose                num_op_rev_tl
##                12.1711694                7.7797112
##                num_actv_rev_tl                num_rev_tl_bal_gt_0
##                4.7089597                4.7071841
##                num_rev_accts                num_bc_tl
##                3.0915479                2.5854790
##                pct_tl_nvr_dlq                percent_bc_gt_75
##                2.3126132                2.1498416
##                propSatisBankcardAccts                num_bc_sats
##                2.0298433                1.0148114
##                pub_rec_bankruptcies                num_il_tl
##                0.5082239                0.4139405
##                bc_util                num_accts_ever_120_pd
##                0.3726118                0.2691630
```

Prune Tree based on cp

```
prn_lcDT <- prune(Wghtd_lcDT, cp=0.0019664)
```

```
rpart.plot::prp(prn_lcDT, type=2, extra=1)
```



Check Performance for Validation Set

```
#Evaluate performance
predVal=predict(prn_lcDT,lcdfVal, type='class')
table(predictValidation = predVal, true=lcdfVal$loan_status)
```

```
##                true
## predictValidation Charged Off Fully Paid
##                Charged Off                245                814
##                Fully Paid                886                6157
```

```
mean(predVal == lcdfVal$loan_status)
```

```
## [1] 0.7901753
```

```
#Confusion table
```

```
confusionMatrix(table(predictValidation = predVal, true=lcdfVal$loan_status))
```

```
## Confusion Matrix and Statistics
```

```
##
```

```
##           true
```

```
## predictValidation Charged Off Fully Paid
```

```
##      Charged Off      245      814
```

```
##      Fully Paid      886     6157
```

```
##
```

```
##           Accuracy : 0.7902
```

```
##           95% CI : (0.7811, 0.799)
```

```
##      No Information Rate : 0.8604
```

```
##      P-Value [Acc > NIR] : 1.00000
```

```
##
```

```
##           Kappa : 0.1026
```

```
##
```

```
##      McNemar's Test P-Value : 0.08507
```

```
##
```

```
##           Sensitivity : 0.21662
```

```
##           Specificity : 0.88323
```

```
##      Pos Pred Value : 0.23135
```

```
##      Neg Pred Value : 0.87420
```

```
##           Prevalence : 0.13960
```

```
##      Detection Rate : 0.03024
```

```
##      Detection Prevalence : 0.13071
```

```
##      Balanced Accuracy : 0.54993
```

```
##
```

```
##      'Positive' Class : Charged Off
```

```
##
```

```
Check the Performance of The Test Set
```

```
#Evaluate performance on the test set
```

```
predTst=predict(prn_lcDT, lcdfTst, type = 'class')
```

```
table(predictTest = predTst, true=lcdfTst$loan_status)
```

```
##           true
```

```
## predictTest  Charged Off Fully Paid
```

```
##   Charged Off      521     1493
```

```
##   Fully Paid     1866     12325
```

```
mean(predTst ==lcdfTst$loan_status)
```

```
## [1] 0.7927183
```

```
confusionMatrix(table(predictTest = predTst, true=lcdfTst$loan_status))
```

```
## Confusion Matrix and Statistics
```

```
##
```

```
##           true
```

```
## predictTest  Charged Off Fully Paid
```

```
##   Charged Off      521     1493
```

```
## Fully Paid      1866      12325
##
##           Accuracy : 0.7927
##           95% CI : (0.7864, 0.7989)
## No Information Rate : 0.8527
## P-Value [Acc > NIR] : 1
##
##           Kappa : 0.1178
##
## McNemar's Test P-Value : 1.376e-10
##
##           Sensitivity : 0.21827
##           Specificity : 0.89195
## Pos Pred Value : 0.25869
## Neg Pred Value : 0.86851
## Prevalence : 0.14730
## Detection Rate : 0.03215
## Detection Prevalence : 0.12428
## Balanced Accuracy : 0.55511
##
## 'Positive' Class : Charged Off
##
```

ROCR For Weighted Rpart Tree

```
scoreTst=predict(prn_lcDT, lcdfTst, type="prob")[, 'Charged Off']
```


```
#apply the prediction function from ROCR to get a prediction object
```

```
rocPredTst = prediction(scoreTst, lcdfTst$loan_status, label.ordering = c('Fully Paid', 'Charged Off'))
```

```
perfROCTst=performance(rocPredTst, "tpr", "fpr")
```

```
plot(perfROCTst)
```

```
abline(0,1)
```



Lifts for Weighted Rpart tree

```
# 'scores' from applying the model to the data
predTrnProb=predict(prn_lcDT, lcdfTrn, type='prob')
head(predTrnProb)

##    Charged Off Fully Paid
## 1    0.2758138  0.7241862
## 2    0.2059635  0.7940365
## 3    0.2758138  0.7241862
## 4    0.2758138  0.7241862
## 5    0.2059635  0.7940365
## 6    0.5798722  0.4201278

#Create a data-frame with only the model scores and the actual class
trnSc <- lcdfTrn %>% select("loan_status")
trnSc$score<-predTrnProb[, 1]

#take a look at trnSc
head(trnSc)

## # A tibble: 6 x 2
##   loan_status score
##   <fct>      <dbl>
## 1 Fully Paid 0.276
## 2 Fully Paid 0.206
## 3 Fully Paid 0.276
## 4 Fully Paid 0.276
## 5 Fully Paid 0.206
## 6 Fully Paid 0.580


#sort by score
trnSc<-trnSc[order(trnSc$score, decreasing=TRUE),]

#generate the cumulative sum of "default" OUTCOME values
trnSc$cumDefault<-cumsum(trnSc$loan_status == "Charged Off")

#first 10 row in trnSc
trnSc[1:10,]

## # A tibble: 10 x 3
##   loan_status score cumDefault
##   <fct>      <dbl>      <int>
## 1 Fully Paid 0.824         0
## 2 Charged Off 0.824         1
## 3 Charged Off 0.824         2
## 4 Charged Off 0.824         3
## 5 Charged Off 0.824         4
## 6 Fully Paid 0.824         4
## 7 Fully Paid 0.824         4
## 8 Fully Paid 0.824         4
## 9 Charged Off 0.824         5
## 10 Charged Off 0.824         6

#Plot the cumDefault values (y-axis) by numCases (x-axis)
plot( trnSc$cumDefault, type = "l", xlab='#cases', ylab='#Charged Off')
abline(0,max(trnSc$cumDefault)/56714, col="blue") #diagonal line
```



Calculate the decile lift table.


```
#Divide the data into 10 for decile lift equal groups
trnSc["bucket"]<- ntile(-trnSc[, "score"], 10)

#group the data by the 'buckets', and obtain summary statistics
dLifts <- trnSc %>% group_by(bucket) %>% summarize(count=n(), numDefaults=sum(loan_status=="Charged Off",
  defRate=numDefaults/count, cumDefRate=cumsum(numDefaults)/cumsum(count),
  lift = cumDefRate/(sum(trnSc$loan_status=="Charged Off")/nrow(trnSc)) )

#look at the table
dLifts
```


```
## # A tibble: 10 x 6
##   bucket count numDefaults defRate cumDefRate lift
## *   <int> <int>      <int>   <dbl>      <dbl> <dbl>
## 1     1    5672      2102  0.371      0.371  2.53
## 2     2    5672      1244  0.219      0.219  1.50
## 3     3    5672      1077  0.190      0.190  1.30
## 4     4    5672       941  0.166      0.166  1.13
## 5     5    5671       635  0.112      0.112  0.764
## 6     6    5671       517  0.0912     0.0912  0.622
## 7     7    5671       433  0.0764     0.0764  0.521
## 8     8    5671       451  0.0795     0.0795  0.543
## 9     9    5671       454  0.0801     0.0801  0.546
## 10    10    5671       455  0.0802     0.0802  0.548
```

```
#various plots,
plot(dLifts$bucket, dLifts$lift, xlab="deciles", ylab="Cumulative Decile Lift", type="l")
```




```
barplot(dLifts$numDefaults, main="numDefaults by decile", xlab="deciles")
```

numDefaults by decile



```
library('lift')
```

```
plotLift(trnSc$score, trnSc$loan_status == "Charged Off")
```



```
#value of lift in the top decile
TopDecileLift(trnSc$score, trnSc$loan_status)
```

```
## [1] NA
```

5 b) continued - C50 Tree

The same model process used for rpart was followed for C50 using the separate training, validation and test sets, and the data was weighted to balance the fully paid and charged off. The training data confusion matrix resulted in 79.71% accuracy and 88.14% specificity. As the training data was at about 80%, the cross validation was run and resulted in 79.1% accuracy and 87.5% specificity. The test model resulted in 78.96% accuracy and 87.9% specificity. The roc curve was mostly convex with fairly poor separation, then the curve was lifted to evaluate performance.

C50 Tree

```
#build a tree model
c5_DT1 <- C5.0(loan_status ~ ., data=lcdfTrn, control=C5.0Control(minCases=50), weights = myweights)
```

Prediction Train, Val, and Test all at once

```
predTrnProb_c5dt1 <- predict(c5_DT1, lcdfTrn, type='class')
predValProb_c5dt1 <- predict(c5_DT1, lcdfVal, type='class')
predTstProb_c5dt1 <- predict(c5_DT1, lcdfTst, type='class')
```

```
#Training
mean(predTrnProb_c5dt1==lcdfTrn$loan_status)
```

```
## [1] 0.7971224
```

```
#Validation
mean(predValProb_c5dt1==lcdfVal$loan_status)
```

```
## [1] 0.7910392
```

```
#Test
```

```
mean(predTstProb_c5dt1==lcdfTst$loan_status)
```

```
## [1] 0.7895711
```

Predictions for Training

```
#method 2
```

```
predTrnProb_c5dt1 <- predict(c5_DT1, lcdfTrn, type='class')
```

```
confusionMatrix(table(predictC50Train = predTrnProb_c5dt1, true=lcdfTrn$loan_status))
```

```
## Confusion Matrix and Statistics
```

```
##
```

```
##           true
```

```
## predictC50Train Charged Off Fully Paid
```

```
##   Charged Off      2545      5742
```

```
##   Fully Paid       5764     42663
```

```
##
```

```
##           Accuracy : 0.7971
```

```
##           95% CI : (0.7938, 0.8004)
```

```
##   No Information Rate : 0.8535
```

```
##   P-Value [Acc > NIR] : 1.0000
```

```
##
```

```
##           Kappa : 0.1879
```

```
##
```

```
## Mcnemar's Test P-Value : 0.8448
```

```
##
```

```
##           Sensitivity : 0.30629
```

```
##           Specificity : 0.88138
```

```
##           Pos Pred Value : 0.30711
```

```
##           Neg Pred Value : 0.88098
```

```
##           Prevalence : 0.14651
```

```
##           Detection Rate : 0.04487
```

```
##   Detection Prevalence : 0.14612
```

```
##           Balanced Accuracy : 0.59384
```

```
##
```

```
##           'Positive' Class : Charged Off
```

```
##
```

```
mean(predTrnProb_c5dt1==lcdfTrn$loan_status)
```

```
## [1] 0.7971224
```

Predictions for Validation

```
predValProb_c5dt1 <- predict(c5_DT1, lcdfVal, type='class')
```

```
confusionMatrix(table(predictC50Validation = predValProb_c5dt1, true=lcdfVal$loan_status))
```

```
## Confusion Matrix and Statistics
```

```
##
```

```
##           true
```

```
## predictC50Validation Charged Off Fully Paid
```

```
##   Charged Off      307      869
```

```
##   Fully Paid      824     6102
```

```
##
```

```
##           Accuracy : 0.791
```

```
##          95% CI : (0.782, 0.7998)
##    No Information Rate : 0.8604
##    P-Value [Acc > NIR] : 1.0000
##
##          Kappa : 0.1444
##
##    McNemar's Test P-Value : 0.2849
##
##          Sensitivity : 0.27144
##          Specificity : 0.87534
##          Pos Pred Value : 0.26105
##          Neg Pred Value : 0.88103
##          Prevalence : 0.13960
##          Detection Rate : 0.03789
##    Detection Prevalence : 0.14515
##          Balanced Accuracy : 0.57339
##
##    'Positive' Class : Charged Off
##
```

Predictions for Test

```
predTstProb_c5dt1 <- predict(c5_DT1, lcdfTst, type='class')
confusionMatrix(table(predictC50Test = predTstProb_c5dt1, true=ldcfTst$loan_status))
```

```
## Confusion Matrix and Statistics
##
##          true
## predictC50Test Charged Off Fully Paid
##    Charged Off          653          1676
##    Fully Paid           1734          12142
##
##          Accuracy : 0.7896
##          95% CI : (0.7832, 0.7958)
##    No Information Rate : 0.8527
##    P-Value [Acc > NIR] : 1.000
##
##          Kappa : 0.1538
##
##    McNemar's Test P-Value : 0.329
##
##          Sensitivity : 0.2736
##          Specificity : 0.8787
##          Pos Pred Value : 0.2804
##          Neg Pred Value : 0.8750
##          Prevalence : 0.1473
##          Detection Rate : 0.0403
##    Detection Prevalence : 0.1437
##          Balanced Accuracy : 0.5761
##
##    'Positive' Class : Charged Off
##
```

ROCR For Weighted C50 Tree


```

#obtain the scores from the model for the class of interest
c5scoreTst=predict(c5_DT1, lcdfTst, type="prob")[, 'Charged Off']

# apply the prediction function from ROCR to get a prediction object
c5rocPredTst = prediction(c5scoreTst, lcdfTst$loan_status, label.ordering = c('Fully Paid', 'Charged Off'))

c5perfROCTst=performance(c5rocPredTst, "tpr", "fpr")
plot(c5perfROCTst)
abline(0,1)

```



Lifts for Weighted Rpart tree

```

#get the 'scores' from applying the model to the data
c5predTrnProb=predict(c5_DT1, lcdfTrn, type='prob')


c5trnSc <- lcdfTrn %>% select("loan_status") # selects the OUTCOME column into trnSc
c5trnSc$score<-c5predTrnProb[, 1]

#sort by score
c5trnSc<-c5trnSc[order(c5trnSc$score, decreasing=TRUE),]

#generate the cumulative sum of "default" OUTCOME values
c5trnSc$cumDefault<-cumsum(c5trnSc$loan_status == "Charged Off")

#Plot the cumDefault values (y-axis) by numCases (x-axis)
plot( c5trnSc$cumDefault, type = "l", xlab='#cases', ylab='#Charged Off')
abline(0,max(c5trnSc$cumDefault)/56714, col="blue") #diagonal line

```



Calculate the decile lift table.


```
#Divide the data into 10 (for decile lift) equal groups
c5trnSc["bucket"]<- ntile(-c5trnSc[, "score"], 10)

#group the data by the 'buckets', and obtain summary statistics
c5dLifts <- c5trnSc %>% group_by(bucket) %>% summarize(count=n(), numDefaults=sum(loan_status=="Charged
defRate=numDefaults/count, cumDefRate=cumsum(numDefaults)/cumsum(count),
lift = cumDefRate/(sum(c5trnSc$loan_status=="Charged Off")/nrow(c5trnSc)) )

#look at the table
c5dLifts
```


```
## # A tibble: 10 x 6
##   bucket count numDefaults defRate cumDefRate lift
## *   <int> <int>      <int>   <dbl>      <dbl> <dbl>
## 1     1   5672      1758  0.310      0.310  2.12
## 2     2   5672      1385  0.244      0.244  1.67
## 3     3   5672      1082  0.191      0.191  1.30
## 4     4   5672       973  0.172      0.172  1.17
## 5     5   5671       726  0.128      0.128  0.874
## 6     6   5671       680  0.120      0.120  0.818
## 7     7   5671       636  0.112      0.112  0.765
## 8     8   5671       465  0.0820     0.0820  0.560
## 9     9   5671       286  0.0504     0.0504  0.344
## 10    10   5671       318  0.0561     0.0561  0.383
```

```
#you can do various plots, for example
plot(c5dLifts$bucket, c5dLifts$lift, xlab="deciles", ylab="Cumulative Decile Lift", type="l")
```





```
barplot(c5dLifts$numDefaults, main="numDefaults by decile", xlab="deciles")
```

numDefaults by decile



```
plotLift(c5trnSc$score, c5trnSc$loan_status == "Charged Off")
```



```
#value of lift in the top decile
TopDecileLift(c5trnSc$score, c5trnSc$loan_status)
```

```
## Warning in TopDecileLift(c5trnSc$score, c5trnSc$loan_status): NAs introduced by
## coercion
## [1] NA
```

5 c) What is your best model?

Rpart is identified as the best decision tree model as it had consistently (although only slightly) higher accuracies for the training, validation and test sets. The size of the tree was fairly small in complexity. Variable importance was determined by the information index. The most important variable in the rpart model was interest rate, followed by subgrade and grade.

6. Develop a Random Forest Model

We decided to use the parameters `min.node.size=1` for classification, `importance='impurity'` to use the Gini index because we are running the model for classification

Due to the imbalance of the data in loan status we used weight in ratio of 5 to 1 for charged off to be compensated. To develop the mode we use the library `ranger`. We decided to use the parameters `min.node.size=1` for classification, `importance='impurity'` to use the Gini index because we are running the model for classification. Also, we used the parameter `case.weights` to balance the data.

We obtained accuracy of 0.85 with the model with the training set, 0.8413 on the test set and 0.8400 on the validation set. The ROC curve was used to evaluate the performance of the model.

```
#Random Forest

library(ranger)

myweights = ifelse(lcdfTrn$loan_status == "Charged Off", 5, 1)
```

```
rgModel1 <- ranger(loan_status ~., data=lcdfTrn, num.trees =200, min.node.size=1, importance='impurity')
#We decided to use the parameters min.node.size=1 for classication, importance='impurity' to use the Gi
```

```
#variable importance
```

```
importance(rgModel1)
```

```
##          loan_amnt          int_rate
##          725.626877          1270.915584
##          installment          grade
##          907.421801          783.816622
##          sub_grade          emp_length
##          994.349131          500.401281
##          home_ownership          annual_inc
##          180.693203          891.388135
##          verification_status          purpose
##          223.112197          285.081721
##          dti          earliest_cr_line
##          1065.277413          899.598588
##          initial_list_status collections_12_mths_ex_med
##          123.760840          32.172330
##          total_rev_hi_lim          acc_open_past_24mths
##          883.265767          573.474678
##          avg_cur_bal          bc_open_to_buy
##          986.236798          928.222655
##          bc_util          chargeoff_within_12_mths
##          903.237585          13.839041
##          delinq_amnt          mo_sin_old_il_acct
##          7.339083          861.358359
##          mo_sin_old_rev_tl_op          mo_sin_rcnt_rev_tl_op
##          924.470823          619.822463
##          mo_sin_rcnt_tl          mort_acc
##          569.573159          340.533635
##          mths_since_recent_bc          mths_since_recent_inq
##          740.265004          628.440252
##          num_accts_ever_120_pd          num_actv_bc_tl
##          212.571457          368.049397
##          num_actv_rev_tl          num_bc_sats
##          416.302106          388.545446
##          num_bc_tl          num_il_tl
##          551.080551          613.152622
##          num_op_rev_tl          num_rev_accts
##          478.100939          621.357695
##          num_rev_tl_bal_gt_0          num_sats
##          416.022612          534.198953
##          num_tl_120dpd_2m          num_tl_30dpd
##          1.379972          9.718062
##          num_tl_90g_dpd_24m          num_tl_op_past_12m
##          79.121872          379.027549
##          pct_tl_nvr_dlq          percent_bc_gt_75
##          527.714033          463.105603
##          pub_rec_bankruptcies          tax_liens
##          111.537678          70.710621
##          tot_hi_cred_lim          total_bal_ex_mort
```

```
##          992.520968          841.593641
##          total_bc_limit total_il_high_credit_limit
##          894.617361          750.163657
##          propSatisBankcardAccts
##          604.877471

rgModel1[["confusion.matrix"]]

##          predicted
## true          Charged Off Fully Paid <NA>
## Charged Off          1579          6727    3
## Fully Paid           1451          46954    0

#pr <- predict (rgModel1, lcdfTst, predict.all = FALSE, proximity = FALSE, type = 'response')

rgModel1[["confusion.matrix"]]

##          predicted
## true          Charged Off Fully Paid <NA>
## Charged Off          1579          6727    3
## Fully Paid           1451          46954    0

#          predicted
#true          Charged Off Fully Paid
# Charged Off          1588          6681
# Fully Paid           1372          47073

(1588+47073)/(1588 +6681+1372+47073) # 0.85

## [1] 0.8580068

#scoreTest
scoresRFTest <- predict(rgModel1, lcdfTst)
#confusion table test data
table(scoresRFTest$predictions,lcdfTst$loan_status)

##          Charged Off Fully Paid
## Charged Off          192          386
## Fully Paid           2195          13432

#          Charged Off Fully Paid
# Charged Off          151          373
#Fully Paid           2198          13483
(151+13483)/(151+13483+373+2198) # 0.8413

## [1] 0.8413453

#scoreVal
scoresRFVal <- predict(rgModel1, lcdfVal)
#confusion table validation
table(scoresRFVal$predictions,lcdfVal$loan_status)

##          Charged Off Fully Paid
## Charged Off          97          203
## Fully Paid           1034          6768
```

```
#           Charged Off Fully Paid
#Charged Off       79       166
#Fully Paid       1130      6727
(79+6727)/(79+6727+1130+166) #0.8400
```

```
## [1] 0.8400395
```

Predictions for Validation


```
predValProb_rgModel1 <- predict(rgModel1, lcdfVal, type='response')
```

Predictions for Test


```
predTstrgModel1 <- predict(rgModel1, lcdfTst, type='response')
```

#ROC

```
library('ROCR')
rgModelROC <- ranger(loan_status ~., data=lcdfTrn, num.trees =200, min.node.size=1, importance='impurity')
scoresRFTest <- predict(rgModelROC, lcdfTst, type="response")
#now apply the prediction function from ROCR to get a prediction object for charge off
rocPredTst <- prediction(scoresRFTest [["predictions"]][,2], lcdfTst$loan_status, label.ordering = c('C'))
#obtain performance using the function from ROCR, then plot
perfROCTst <- performance(rocPredTst, "tpr", "fpr")
plot(perfROCTst)
```



```
#now apply the prediction function from ROCR to get a prediction object for fully paid
rocPredTst <- prediction(scoresRFTest [["predictions"]][,1], lcdfTst$loan_status, label.ordering = c('F'))
#obtain performance using the function from ROCR, then plot
perfROCTst <- performance(rocPredTst, "tpr", "fpr")
plot(perfROCTst)
```



Loans Analysis

```
library(lubridate)
```

```
#loans by grade
```

```
lcdf2 %>% group_by(grade) %>% summarise(nLoans=n(), defaults=sum(loan_status=="Charged Off"),  
avgInterest= mean(int_rate), stdInterest=sd(int_rate), avgLoanAMt=mean(loan_amnt), avgPmnt=mean(total_pymnt))
```

```
## # A tibble: 7 x 7
```

##	grade	nLoans	defaults	avgInterest	stdInterest	avgLoanAMt	avgPmnt
## 1	A	20402	1108	7.25	0.796	14146.	15188.
## 2	B	23398	2682	10.7	1.22	12459.	13549.
## 3	C	22577	4116	13.7	0.850	11466.	12363.
## 4	D	10802	2647	16.5	0.895	12150.	12957.
## 5	E	3191	1045	19.8	1.10	12558.	13079.
## 6	F	560	191	24.1	0.798	10169.	10588.
## 7	G	91	38	25.8	0.0593	12509.	13576.

```
#calculate the annualized percentage return
```

```
lcdf2$annRet <- ((lcdf2$total_pymnt -lcdf2$funded_amnt)/lcdf2$funded_amnt)*(12/36)*100
```

```
#summarize by grade
```

```
lcdf2 %>% group_by(grade) %>% summarise(nLoans=n(), defaults=sum(loan_status=="Charged Off"), avgInterest=mean(int_rate),  
minRet=min(annRet), maxRet=max(annRet))
```

```
## # A tibble: 7 x 11
```

##	grade	nLoans	defaults	avgInterest	stdInterest	avgLoanAMt	avgPmnt	avgRet	stdRet
## 1	A	20402	1108	7.25	0.796	14146.	15188.	2.39	3.88
## 2	B	23398	2682	10.7	1.22	12459.	13549.	2.85	5.95
## 3	C	22577	4116	13.7	0.850	11466.	12363.	2.65	8.00
## 4	D	10802	2647	16.5	0.895	12150.	12957.	2.32	9.76
## 5	E	3191	1045	19.8	1.10	12558.	13079.	1.40	11.7

```
## 6 F      560      191      24.1      0.798      10169.  10588.   2.55  12.7
## 7 G       91       38      25.8      0.0593    12509.  13576.   2.09  13.4
## # ... with 2 more variables: minRet <dbl>, maxRet <dbl>

#Some loans are paid back early - find out the actual loan term in months
lcdf2$last_pymnt_d<-paste(lcdf2$last_pymnt_d, "-01", sep = "")
lcdf2$last_pymnt_d<-parse_date_time(lcdf2$last_pymnt_d, "mYd")

# getting actual term
lcdf2 $actualTerm <- ifelse(lcdf2$loan_status=="Fully Paid", as.duration(lcdf2$issue_d %--% lcdf2$last_pymnt_d)/365, 0)

#Then, considering this actual term, the actual annual return is
lcdf2$actualReturn <- ifelse(lcdf2$actualTerm>0, ((lcdf2$total_pymnt - lcdf2$funded_amnt)/lcdf2$funded_amnt)^(365/lcdf2$actualTerm)-1, 0)

#loan performance by grade
lcdf2 %>% group_by(grade) %>% summarise(nLoans=n(), defaults=sum(loan_status=="Charged Off"), defaultRate=
avgInterest= mean(int_rate), avgLoanAmt=mean(loan_amnt), avgRet=mean(annRet), avgActualRet=mean(actualReturn),
avgActualTerm=mean(actualTerm), minActualRet=min(actualReturn)*100, maxActualRet=max(actualReturn)*100)

## # A tibble: 7 x 11
##   grade nLoans defaults defaultRate avgInterest avgLoanAmt avgRet avgActualRet
## * <fct> <int> <int> <dbl> <dbl> <dbl> <dbl> <dbl>
## 1 A      20402    1108    0.0543    7.25    14146.   2.39    3.94
## 2 B      23398    2682    0.115    10.7    12459.   2.85    5.22
## 3 C      22577    4116    0.182    13.7    11466.   2.65    5.73
## 4 D      10802    2647    0.245    16.5    12150.   2.32    5.89
## 5 E       3191    1045    0.327    19.8    12558.   1.40    5.45
## 6 F       560     191    0.341    24.1    10169.   2.55    7.29
## 7 G        91      38    0.418    25.8    12509.   2.09    6.33
## # ... with 3 more variables: avgActualTerm <dbl>, minActualRet <dbl>,
## #   maxActualRet <dbl>

#loan performance by grade and loan status
lcdf2 %>% group_by(grade, loan_status) %>% summarise(nLoans=n(), defaults=sum(loan_status=="Charged Off"),
avgInterest= mean(int_rate), avgLoanAmt=mean(loan_amnt), avgRet=mean(annRet), avgActualRet=mean(actualReturn),
avgActualTerm=mean(actualTerm), minActualRet=min(actualReturn), maxActualRet=max(actualReturn))

## # A tibble: 14 x 12
## # Groups:   grade [7]
##   grade loan_status nLoans defaults defaultRate avgInterest avgLoanAmt avgRet
##   <fct> <fct> <int> <int> <dbl> <dbl> <dbl> <dbl>
## 1 A      Charged Off  1108    1108    1      7.47    13438. -11.2
## 2 A      Fully Paid  19294     0      0      7.24    14187.   3.18
## 3 B      Charged Off  2682    2682    1     10.9    12261. -11.0
## 4 B      Fully Paid  20716     0      0     10.7    12484.   4.64
## 5 C      Charged Off  4116    4116    1     13.7    11652. -11.6
## 6 C      Fully Paid  18461     0      0     13.7    11424.   5.83
## 7 D      Charged Off  2647    2647    1     16.6    12495. -12.2
## 8 D      Fully Paid  8155     0      0     16.5    12038.   7.04
## 9 E      Charged Off  1045    1045    1     19.8    12449. -12.8
## 10 E     Fully Paid  2146     0      0     19.7    12611.   8.31
## 11 F     Charged Off   191     191    1     24.2    11158. -12.5
## 12 F     Fully Paid   369     0      0     24.1     9657.  10.4
## 13 G     Charged Off    38     38     1     25.8    11250 -11.0
## 14 G     Fully Paid    53     0      0     25.8    13412.  11.5
## # ... with 4 more variables: avgActualRet <dbl>, avgActualTerm <dbl>,
```

```
## # minActualRet <dbl>, maxActualRet <dbl>
#profitValue based on
lcdf2 %>% group_by(loan_status) %>% summarise(avgInt=mean(int_rate),avgActInt = mean(actualReturn))

## # A tibble: 2 x 3
##   loan_status avgInt avgActInt
## * <fct>      <dbl>    <dbl>
## 1 Charged Off  13.9    -0.117
## 2 Fully Paid   11.6     0.0803

PROFITVAL <- 24 #profit
COSTVAL <- -35 # loss

Avg = 8.03 * 2.1 + 2*0.9

#Performance
scoreTstRF2 <- predict(rgModel1,lcdfTst, type="response")["Fully Paid"]

prPerfRF2 <- data.frame(scoreTstRF2)
#prPerfRF2 <- cbind(prPerfRF2, status=lcdfTst$loan_status)
#prPerfRF2 <- prPerfRF2[order(-scoreTstRF2) ,] #sort in desc order of prob(fully_paid) prPerfRF$profit
#max(prPerfRF$cumProfit) prPerfRF$cumProfit[which.max(prPerfRF$cumProfit)]
```

7. (a). Evaluate the Loans for Investment Decisions

rpart model

	true	
predictTest	Charged Off	Fully Paid
Charged Off	521	1493
Fully Paid	1866	12325

$$8.03 \times 2.1 + (.9) \times 2 = 17.04 \quad 12325 = 210,054.98, \quad (-11.7) \times 3 = -35.1 \times 1866 = 65,496.6.$$

$$210,054.98 - 65,496.6 = 144,558.36 \text{ profit for the rpart model}$$

C50 model

	true	
predictC50Test	Charged Off	Fully Paid
Charged Off	653	1676
Fully Paid	1734	12142

$$8.03 \times 2.1 + (.9) \times 2 = 17.04 \quad 12142 = 206,936.11, \quad (-11.7) \times 3 = -35.1 \times 1734 = 60,863.4$$

$$206,936.11 - 60,863.4 = 146,072.71$$

When doing these profit evaluations, we used the average returns on both “Charged off” and “Fully Paid” loans. We then used our models to look at the predictions and how we fared. In the matrix where we predicted it would be Fully paid and it was fully paid we multiplied that by the average amount for a fully paid loan to get profit. Then for the loans that we thought would be fully paid and end up be Charged off we multiplied that amount by average return to get what we would have lost be the prediction model. The other instances we would not have invested if we thought it would be charged off and it ended up being Fully paid, and the predictions we got correct for Charged Off we would not have invested in either so not lost money.

7 (b).

if you look at the data in a descending order by the probability of becoming Fully Paid, you can see certain points where the drop offs take place and the percentage of defaults greatly increases. We chose the score to

cut off at .589 which has a 18.7% default rate. We felt that after this point the score went below .500 and that the default percentage got closer to 25% which can become riskier. We also looked at the amount of loans that were in each score and most of them were before this cutoff, this shows that a vast majority of these loans scored well in our model.

The advantageous part about using the model like this is that it's weighted with so many high scores showing a default rate at only 7.96%. Almost half of all loans are in that score zone. When comparing using a model like this to invest in safe cd's the risk is the greatest factor here because these safe cds are guaranteed to be paid out. Safe cds provide the 2% interest every year. On 100 dollars that will turn in \$106 by the end of year 3. Using our model is more profitable than investing in safe cds because if you have the same amount of loans as from our previous models you only end up with 97,230 in profit. You get this from taking the average return of 2 dollars for 3 years and then multiplying that by the number of loans at 16205. This is still a decent return but if you are doing a lot of loans at once you have the ability to spread the risk out more and find potentially more profitable loans. That 97,230 compared to above models at over \$140,000 is a drastic difference over a short period of time.

```
#get the 'scores' from applying the model to the data
```

```
predTrnProb2=predict(prn_lcDT, lcdfTrn, type='prob')
```

```
trnSc2 <- lcdfTrn %>% select("loan_status") # selects the OUTCOME column into trnSc
```

```
trnSc2$score<-predTrnProb2[, 2] #add a column named 'Score' with prob(default) values in the first col
```

```
#sort by score
```

```
trnSc2<-trnSc2[order(trnSc2$score, decreasing=TRUE),]
```

```
trnSc2[1:50,]
```

```
## # A tibble: 50 x 2
##   loan_status score
##   <fct>      <dbl>
## 1 Fully Paid  0.794
## 2 Fully Paid  0.794
## 3 Fully Paid  0.794
## 4 Fully Paid  0.794
## 5 Charged Off 0.794
## 6 Fully Paid  0.794
## 7 Fully Paid  0.794
## 8 Fully Paid  0.794
## 9 Fully Paid  0.794
## 10 Fully Paid 0.794
## # ... with 40 more rows
```

```
trnSc2 %>% group_by(score, loan_status) %>% summarise(nloans = n())
```

```
## `summarise()` has grouped output by 'score'. You can override using the `.groups` argument.
```

```
## # A tibble: 32 x 3
## # Groups:   score [16]
##   score loan_status nloans
##   <dbl> <fct>      <int>
## 1 0.176 Charged Off    56
## 2 0.176 Fully Paid    36
## 3 0.327 Charged Off   123
## 4 0.327 Fully Paid   179
## 5 0.358 Charged Off   490
## 6 0.358 Fully Paid   819
```

```
## 7 0.362 Charged Off 1117
## 8 0.362 Fully Paid 1904
## 9 0.407 Charged Off 286
## 10 0.407 Fully Paid 589
## # ... with 22 more rows

trnSc2 %>% group_by(score) %>% summarise(nLoans=n(), defaults=sum(loan_status=="Charged Off")) %>% mutate(

## # A tibble: 16 x 4
##   score nLoans defaults prctCharged_off
## * <dbl> <int> <int> <dbl>
## 1 0.176 92 56 60.9
## 2 0.327 302 123 40.7
## 3 0.358 1309 490 37.4
## 4 0.362 3021 1117 37.0
## 5 0.407 875 286 32.7
## 6 0.420 768 242 31.5
## 7 0.452 767 221 28.8
## 8 0.547 804 174 21.6
## 9 0.589 9997 1884 18.8
## 10 0.592 615 115 18.7
## 11 0.601 469 85 18.1
## 12 0.618 1254 214 17.1
## 13 0.653 2303 346 15.0
## 14 0.724 6995 788 11.3
## 15 0.731 265 29 10.9
## 16 0.794 26878 2139 7.96
```

xgboost model

Additional models to consider - develop boosted tree models (using either gbm or XGBoost). Explain how you experiment with parameters, how performance varies, which parameter setting you use for the ‘best’ model.

Model performance should be evaluated through use of same set of criteria as for the other models - confusion matrix based, ROC analyses and AUC, cost-based performance.

Provide a table with comparative evaluation of all the best models from each methods; show their ROC curves in a combined plot. Also provide profit-curves and ‘best’ profit’ and associated cutoff. At this cutoff, what are the accuracy values for the different models?

As a first step we prepared the data to be used in xgboost. For this the data was converted to numeric with the use of the caret library, all variables were converted to dummy vars except the dependent variable loan status. Then we created a new dataset called dxlcdf with the function predict. For loan status we converted as a dummy variable and kept the level charged off. Then we create the training, test and validation sets.

Next we took care of the unbalanced data for this we calculate the $\frac{\text{sqr}(\text{sum}(\text{negative instances}))}{\text{sum}(\text{positive instances})}$ we apply this number later to the parameter `scale_pos_weight`. Before doing this the model would run the 500 rounds but after balancing the weights we reduce to a best iteration of 49 with an accuracy of 0.85 We then calculate the xgboost for the validation data and get the best iteration on number 3.

```
library(caret)
library(xgboost)
```

```
## Warning: package 'xgboost' was built under R version 3.6.2
```

```
##
```

```

## Attaching package: 'xgboost'

## The following object is masked from 'package:dplyr':
##
##      slice

# Using dummyVars function in the 'caret' package to convert factor variables to dummy-variables.
fdum<-dummyVars(~.,data=lcdf %>% select(-loan_status))

#replacing the dummy variables in the dataset
dxlcdf <- predict(fdum, lcdf)

#checking levels to know how is composed loan status
levels(lcdf$loan_status)

## [1] "Charged Off" "Fully Paid"

#"Fully Paid" "Charged Off"
#converting loan status to dummy variables
dylcdf <- class2ind(lcdf$loan_status, drop2nd = FALSE)
# we decided we want to keep charged off
fplcdf <- dylcdf [ , 2]

#Training, test subsets
dxlcdfTrn <- dxlcdf[indicesTraining,]
colcdfTrn <- fplcdf[indicesTraining]
dxlcdfTst <- dxlcdf[indicesTest,]
colcdfTst <- fplcdf[indicesTest]
dxlcdfVal <- dxlcdf[indicesValidation,]
colcdfVal <- fplcdf[indicesValidation]

#calculating the weights of the subsets
sum(dxlcdfTrn==1)

## [1] 547717
sum(dxlcdfTrn==0)

## [1] 38153380
sqrt(sum(dxlcdfTrn==0) / sum(dxlcdfTrn==1)) #8.823873

## [1] 8.346193
sum(dxlcdfTst==1)

## [1] 156661
sum(dxlcdfTst==0)

## [1] 10901452
sqrt(sum(dxlcdfTst==1) / sum(dxlcdfTst==0))

## [1] 0.1198777
sum(dxlcdfVal==1)

## [1] 78253

```

```

sum(dx1cdfVal==0)

## [1] 5450468
sqrt(sum(dx1cdfVal==1) / sum(dx1cdfVal==0))

## [1] 0.1198212

#Creating of xgb.DMatrix
dxTrn <- xgb.DMatrix(subset(dx1cdfTrn), label=colcdfTrn)
dxTst <- xgb.DMatrix(subset(dx1cdfTst), label=colcdfTst)
dxVal <- xgb.DMatrix(subset(dx1cdfVal), label=colcdfVal)

## Process for training and test

#we use the xgbWatchlist to watch the progress of learning thru performance on these datasets
xgbWatchlist <- list(train = dxTrn, eval = dxTst)

#This is the list of parameters for the xgboost model development functions wich are going to use first
xgbParam <- list (
  max_depth = 5, eta = 0.01, scale_pos_weight = 8.82,
  objective = "binary:logistic", eval_metric="error", eval_metric = "auc")

#confusion matrix
table(pred=as.numeric(xpredTrg>0.5), act=colcdfTrn)

##      act
## pred    0    1
##    0    19    0
##    1 8290 48405

#      act
# pred    0    1
#    0    19    0
#    1 8290 48405
(19 + 48405) / (19+8290+48405)


## [1] 0.853828

# = 0.85 accuracy

#ROC, AUC performance
xpredTst<-predict(xgb_lsM1, dxTst)

pred_xgb_lsM1<-prediction(xpredTst, lcdfTst$loan_status,
  label.ordering = c("Charged Off", ("Fully Paid")))
aucPerf_xgb_lsM1<-performance(pred_xgb_lsM1, "tpr", "fpr")
plot(aucPerf_xgb_lsM1)
abline(a=0, b= 1)

```



Using cross-validation on training dataset to determine best model

```
xgbParamGrid
```

```
##   max_depth  eta bestTree bestPerf
## 1         2 0.001      51 0.670853
## 2         5 0.001      73 0.672626
## 3         2 0.010      15 0.665226
## 4         5 0.010      38 0.673593
## 5         2 0.100     147 0.685824
## 6         5 0.100     100 0.684785
```

```
# max_depth eta bestTree bestPerf
# 2 0.001   28 0.668216    <--
# 5 0.001    6 0.671976
# 2 0.010   10 0.670330
# 5 0.010   26 0.674678
# 2 0.100   84 0.682619
# 5 0.100  101 0.684982
```

```
#Best parameters
```

```
xgbParam_Best <- list (booster = "gbtree", objective = "binary:logistic", min_child_weight=1, colsample
```

```
# XGBOOST running the model with the best parameters found with the for loop
```

```
xgb_lsM2 <- xgb.train(xgbParam_Best, dxTrn, nrounds = xgb_tune$best_iteration)
```

```
#XGBOOST evaluation of the model
```

```
#Using the predicting function to get the scores in the training data set
```

```
xpredTrn<-predict(xgb_lsM2, dxTrn)
```

```
#Using the predicting function to get the scores in the test data set
```

```
xpredTst<-predict(xgb_lsM2, dxTst)
```

```
#confusion matrix
```


```
table(pred=as.numeric(xpredTst>0.5), act=colcdfTst)
```

```
##      act
## pred    0    1
##      1 2387 13818
```

```
#ROC, AUC performance
```

```
pred_xgb_lsM2<-prediction(xpredTst, lcdfTst$loan_status,
label.ordering = c("Charged Off", ("Fully Paid")))
```

```
aucPerf_xgb_lsM2<-performance(pred_xgb_lsM2, "tpr", "fpr")
plot(aucPerf_xgb_lsM2)
abline(a=0, b= 1)
```



```
#####
```

```
## Process for test and validation
```

```
#we can watch the progress of learning thru performance on these datasets
```

```
xgbWatchlistVal <- list(train = dxTst, eval = dxVal)
```

```
#list of parameters for the xgboost model development functions
```

```
xgbParam <- list (
max_depth = 5, eta = 0.01, scale_pos_weight = 8.82,
objective = "binary:logistic", eval_metric="error", eval_metric = "auc")
```

```
xgb_lsM2 <- xgb.train(xgbParam, dxTst, nrounds = 500, xgbWatchlistVal, early_stopping_rounds = 10 )
```

```
## [1] train-error:0.146560 train-auc:0.652217 eval-error:0.140212 eval-auc:0.643278
```

```
## Multiple eval metrics are present. Will use eval_auc for early stopping.
```

```
## Will train until eval_auc hasn't improved in 10 rounds.
```

```
##
```

```
## [2] train-error:0.146745 train-auc:0.667170 eval-error:0.140089 eval-auc:0.660555
```

```
## [3] train-error:0.146560 train-auc:0.667212 eval-error:0.139595 eval-auc:0.660537
```

```
## [4] train-error:0.146930 train-auc:0.667212 eval-error:0.139595 eval-auc:0.660538
```

```
## [5] train-error:0.146560 train-auc:0.667409 eval-error:0.139595 eval-auc:0.660500
## [6] train-error:0.146683 train-auc:0.667410 eval-error:0.139348 eval-auc:0.660501
## [7] train-error:0.146683 train-auc:0.668102 eval-error:0.139348 eval-auc:0.659921
## [8] train-error:0.146683 train-auc:0.671860 eval-error:0.139348 eval-auc:0.662652
## [9] train-error:0.146683 train-auc:0.671897 eval-error:0.139348 eval-auc:0.662473
## [10] train-error:0.146683 train-auc:0.671858 eval-error:0.139348 eval-auc:0.662522
## [11] train-error:0.146683 train-auc:0.671863 eval-error:0.139348 eval-auc:0.662488
## [12] train-error:0.146683 train-auc:0.675830 eval-error:0.139348 eval-auc:0.663956
## [13] train-error:0.146683 train-auc:0.675819 eval-error:0.139348 eval-auc:0.663961
## [14] train-error:0.146683 train-auc:0.675839 eval-error:0.139348 eval-auc:0.663951
## [15] train-error:0.146683 train-auc:0.675839 eval-error:0.139348 eval-auc:0.663951
## [16] train-error:0.146868 train-auc:0.676188 eval-error:0.139472 eval-auc:0.663943
## [17] train-error:0.146930 train-auc:0.676985 eval-error:0.139595 eval-auc:0.664118
## [18] train-error:0.146930 train-auc:0.676976 eval-error:0.139595 eval-auc:0.664119
## [19] train-error:0.146930 train-auc:0.677552 eval-error:0.139595 eval-auc:0.663914
## [20] train-error:0.146930 train-auc:0.677529 eval-error:0.139595 eval-auc:0.663927
## [21] train-error:0.146868 train-auc:0.677523 eval-error:0.139472 eval-auc:0.663936
## [22] train-error:0.146930 train-auc:0.678568 eval-error:0.139595 eval-auc:0.663667
## [23] train-error:0.146868 train-auc:0.678539 eval-error:0.139472 eval-auc:0.663680
## [24] train-error:0.146930 train-auc:0.678816 eval-error:0.139595 eval-auc:0.663697
## [25] train-error:0.146930 train-auc:0.679338 eval-error:0.139595 eval-auc:0.663410
## [26] train-error:0.146930 train-auc:0.679868 eval-error:0.139595 eval-auc:0.665170
## [27] train-error:0.146930 train-auc:0.681189 eval-error:0.139595 eval-auc:0.665171
## [28] train-error:0.146930 train-auc:0.681220 eval-error:0.139595 eval-auc:0.665179
## [29] train-error:0.146930 train-auc:0.681123 eval-error:0.139595 eval-auc:0.665193
## [30] train-error:0.146930 train-auc:0.681233 eval-error:0.139595 eval-auc:0.665020
## [31] train-error:0.146930 train-auc:0.681358 eval-error:0.139595 eval-auc:0.664875
## [32] train-error:0.146930 train-auc:0.681350 eval-error:0.139595 eval-auc:0.664754
## [33] train-error:0.146930 train-auc:0.681319 eval-error:0.139595 eval-auc:0.664805
## [34] train-error:0.146930 train-auc:0.681364 eval-error:0.139595 eval-auc:0.664789
## [35] train-error:0.146930 train-auc:0.681366 eval-error:0.139595 eval-auc:0.664799
## [36] train-error:0.146930 train-auc:0.681374 eval-error:0.139595 eval-auc:0.664749
## [37] train-error:0.146930 train-auc:0.684078 eval-error:0.139595 eval-auc:0.664669
## [38] train-error:0.146930 train-auc:0.684117 eval-error:0.139595 eval-auc:0.664630
## [39] train-error:0.146930 train-auc:0.684098 eval-error:0.139595 eval-auc:0.664640
## Stopping. Best iteration:
## [29] train-error:0.146930 train-auc:0.681123 eval-error:0.139595 eval-auc:0.665193
```

```
xpredTrg2<-predict(xgb_lsM2, dxVal)
```

```
#confusion matrix
```

```
table(pred=as.numeric(xpredTrg2>0.5), act=colcdfVal)
```

```
##      act
## pred   0    1
##      1 1131 6971
```


```
#ROC, AUC performance
```

```
xpredVal<-predict(xgb_lsM1, dxVal)
```

```
pred_xgb_lsM2<-prediction(xpredTst, lcdfTst$loan_status,
label.ordering = c("Charged Off", ("Fully Paid")))
```

```
aucPerf_xgb_lsM2<-performance(pred_xgb_lsM2, "tpr", "fpr")
```


```
plot(aucPerf_xgb_lsM2)
abline(a=0, b= 1)
```



Plotting Lines For ROC Curves

```
plot(perfROCTst, col="red")
plot(c5perfROCTst, col="blue", add= TRUE)
plot(perfROCTst, col="green", add= TRUE)
plot(aucPerf_xgb_lsM2, col="purple", add=TRUE)
abline(a=0, b= 1, col="black")
```

```
legend(0.8, 0.8, legend=c("rpart", "C5.0", "Ranger", "XgBoost"),
      col=c("red", "blue", "green", "purple"), lty=1:2, cex=0.8)
```

““